

Kill-Chain Canaries: Stage-Level Tracking of Prompt Injection Across Attack Surfaces and Five Production LLMs

Haochuan Kevin Wang¹ Zechen Zhang²

¹Massachusetts Institute of Technology hcw@mit.edu

²University of Chicago zechenz@uchicago.edu

Keywords: prompt injection, LLM security, agent safety, multi-agent systems, red-teaming




Figure 1: Same attack success, different kill chains. PDF relay, visible-text payload (pdf_append), no defense, the same model as Agent A and Agent B; $n = 3$ runs per row (Table 4). (a) Outcome-only scoring gives both models 0/3. (b) The canary token shows where the runs differ: with Claude Haiku 4.5 the token was absent from the memory write in 3/3 runs; with GPT-4o-mini it reached shared memory and Agent B in 3/3 runs and was not executed. Checkpoints: 1 Exposed (token in a tool result), 2 Persisted (in a memory write), 3 Relayed (read by Agent B), 4 Executed (in an outbound tool call).

Abstract

Multi-agent LLM systems are entering production—processing documents, managing workflows, acting **now read documents, web pages and tool results** on behalf of users—yet their resilience users, yet their resistance to prompt injection is still evaluated with a single binary usually reported as one number: did the attack succeed? This leaves architects without the diagnostic information needed to harden real pipelines. We introduce a kill-chain canary methodology that tracks a cryptographic token through four stages (EXPOSED method that plants a unique token in every in-

jected payload and records the furthest of four stages it reaches (Exposed → **PERSISTED** Persisted → **RELAYED** Relayed → **EXECUTED** Executed), across 950 runs, five frontier production LLMs, six attack surfaces, and five defense conditions. The results reframe prompt injection as a pipeline-architecture problem: every model is fully exposed, yet outcomes diverge downstream—Claude blocks all injections at memory-write Exposure was 100% among runs that called the tool; the outcomes differ downstream. Claude Haiku 4.5 and Claude Sonnet 4.5 executed none of their 164 text-surface attacks, and in the text relay the canary token never ap-

peared in a memory write (0/164 ASR, 40); GPT-4o-mini propagates at executed 53% , and DeepSeek exhibits 0%/100% across surfaces from the same model. Three findings matter for deployment: Of its attacks. Four findings follow. (1) write-node placement is the highest-leverage safety decision—routing writes through a verified model eliminates propagation; A Claude writer kept the canary token out of shared memory in every relay run we report; one cross-model pairing (Claude writer, GPT-4o-mini reader, $n = 3$) is consistent with this protecting the reader, and other pairings were not tested. (2) all four defenses fail on at least one surface due to channel mismatch alone, no adversarial adaptation required; As readers, the Claude models executed 0/40 raw pre-seeded injections, but Claude Haiku 4.5 executed 2/3 injections relayed by GPT-4o-mini; whether relayed injections are harder to refuse than raw ones is an open question. (3) invisible whitefont PDF payloads match or exceed visible-text ASR, meaning rendered-layer screening is insufficient. These dynamics apply directly to production: institutional investors and financial firms already run NLP pipelines over earnings calls, SEC filings, and analyst reports—the document-ingestion workflows now migrating to LLM agents. Code , run logs , and tooling DeepSeek Chat went from 0/24 on pre-seeded memory to 8/8 on tool results, scenarios that also differ in task and payload format; white-text PDF payloads, invisible on the rendered page, succeeded at least as often as visible ones. (4) pi_detector and write_filter failed on channels they do not inspect, spotlighting failed on content it wraps, and write_filter blocked the PDF relay but not the text relay, a difference we cannot explain. Code and run logs are publicly released.¹

1 Introduction

Prompt injection—adversarial instructions embedded in data processed by an LLM agent—is now a central attack class for agentic AI deployments [2, 3]. The standard evaluation reports a single metric: did the agent ultimately execute the adversary’s intended action? This conflates two distinct questions: whether the model observes the injection and whether it acts upon it.

Consider a concrete chain: an agent calls `get_webpage()` `get_webpage`, receives poisoned content, calls `write_memory()` `write_memory` to summarize, and is queried by a second agent via `read_memory()` `read_memory`. A 0% ASR for the second agent could mean the injection was stripped at `write_memory` never reached the memory write, or it could mean it survived but the second agent refused to execute. These are architecturally different outcomes: in the first the first implies the summarizer is a decontamination stage; the second implies summarizer dropped the injected text; in the second the terminal agent holds the defense held. Outcome-only

measurement cannot distinguish them; Figure 1 shows two such pipelines with the same 0/3 attack success.

We address this with a kill-chain canary methodology. Every injected payload contains a unique `SECRET-[A-F0-9]{8}` token tracked at four discrete stages by our `PropagationLogger`: every injected payload carries a unique canary token, and our `PropagationLogger` records the furthest of four stages that token reaches. We test five frontier models production LLMs across six attack surfaces (web text, page, pre-seeded memory, tool stream, PDF, invisible PDF result, visible PDF text, white PDF text, and audio transcript) and five defense conditions in `agent_bench` `agent_bench`, a custom multi-agent evaluation harness . A Phase 3 multi-modal extension (Figure 2). A PDF relay experiment evaluates a three-boundary kill chain (document extraction → memory write → agent delegation), with same-model pairs and cross-model relay pairs where pairs in which Agent A and Agent B are come from different model families.

Contributions: (1) Stage-level evaluation framework—A stage-level evaluation framework: a kill-chain canary methodology [1] that attributes defense effectiveness to specific pipeline stages rather than final outcomes records the stage at which each injection stops rather than only the final outcome. Applying it gives four findings. (2) Empirical defense localization — evidence across six injection surfaces and five models that the safety gap concentrates at the summarization write stage, not at context exposure or execution Writer position: a Claude writer kept the canary token out of shared memory in every relay run we report; one cross-model pairing (Claude writer, GPT-4o-mini reader, $n = 3$) is consistent with this protecting the reader, and other pairings were not tested. (3) Cross-surface vulnerability analysis — the first systematic multi-surface evaluation showing a single model’s ASR spans Reader position: as readers, the Claude models executed 0%–100% depending on injection channel, demonstrating that surface coverage determines apparent safety posture/40 raw pre-seeded injections, but Claude Haiku 4.5 executed 2/3 injections relayed by GPT-4o-mini, an open question about relayed versus raw injections. (4) Write-vs-read relay asymmetry — pilot evidence that Claude’s write-stage defense does not extend to read-stage resistance in heterogeneous multi-agent pipelines, suggesting that relay node position (not identity alone) determines downstream safety Channel dependence: DeepSeek Chat went from 0/24 on pre-seeded memory to 8/8 on tool results, scenarios that also differ in task and payload format; white-text PDF payloads succeeded at least as often as visible ones. (5) Defense coverage: pi_detector and write_filter failed on channels they do not inspect, spotlighting failed on content it wraps, and write_filter blocked the PDF relay but not the text relay.

Practical relevance. Many institutional investors and quantitative funds already deploy multi-step The same document-ingestion pattern arises in financial NLP pipelines over unstructured financial text—earnings-call earnings-call transcripts, SEC filings , analyst reports . As these workflows migrate to LLM-agent architectures, every document-ingestion stage becomes a potential injection surface. Our results provide concrete guidance—write-node placement and surface-aware defense composition—for securing these document-driven agent

¹https://github.com/KevinChunye/prompt_injection

deployments and analyst reports as they move to LLM agents.

2 Related Work

AgentDojo [4] provides 97 tasks with 629 injections and reports joint utility-ASR metrics across four environments; it does not decompose by pipeline stage. Our utility-ASR scatter (Figure ??) is directly comparable to their result. InjecAgent [5] evaluates 1,054 indirect-injection cases with a single outcome metric. Prompt Infection [6] demonstrates LLM-to-LLM self-replicating attacks; we identify the exact stage (write_memory summarization) where replication is blocked record the stage (the write_memory summarization) at which the canary token stops. Zombie Agents [?] shows summarization agents persist injections; we replicate and extend to multi-agent relay with quantified stage fractions [7] shows that an injection an agent writes into its long-term memory during a benign session can persist and later trigger unauthorized tool use; we measure that write step directly with a canary token in a two-agent relay. Nasr et al. [?] show adaptive attacks achieve >90% ASR against 12 [8] show that adaptive attacks bypass all 12 defenses they evaluate, most with attack success above 90%. In our evaluation, lightweight defenses failed against non-adaptive attacks achieve the same result via surface mismatch—a structurally distinct failure mode that does not require adaptive adversaries: pi_detector and write_filter on channels they do not inspect, spotlighting on content it wraps. AgentWatcher [11] proposes rule-based, causally attributed prompt injection monitoring; our stage-level canary instrumentation provides a complementary empirical basis for where each model’s defense activates within the pipeline. Xiang et al. [?] [12] survey system-level defenses for indirect prompt injection; our surface-mismatch channel-level results provide quantitative evidence for their observation that defense effectiveness is deployment-context-dependent. Ding et al. [?] study adversarial prompt injection on [13] study covert visual prompt injection against commercial multimodal LLMs; our Phase 3 PDF cross-modal relay experiments extend this PDF relay experiment extends document injection to multi-agent relay chains with per-stage kill-chain tracking. Zhang et al. [14] evaluate privilege usage of LLM agents on with real-world tools; our permission_esc permission_esc scenario finds near-zero ASR (2/132) for the same payload design, suggesting privilege escalation requires more sophisticated payload construction than simple text instructions. Wang et al. [15] study real-time monitoring for reasoning vulnerabilities; their post-hoc detection framing aligns with our finding that objective drift is forensic rather than preventive (leave-one-scenario-out AUC 0.39–0.57). Lynch [16] demonstrates documents the persistent vulnerability of aligned AI systems to adversarial inputs; our finding that Claude’s resistance is localized to the write stage (not the read stage) nuances this—alignment may provide write-time filtering without guaranteeing read-time resistance in heterogeneous pipelines reader results (the Claude models executed 0/40 raw pre-seeded injections, Claude Haiku 4.5 executed 2/3 relayed ones) raise, but do not

settle, the question of whether such resistance differs between writing and reading positions in a pipeline.

3 Benchmark Design

3.1 System Architecture

agent_bench is Figure 2 illustrates our harness, agent_bench, a minimal multi-agent evaluation harness (~setup of about 600 lines of Python) illustrated in Figure 2. Four components are central to the experimental design. MemoryStore provides a key-value store with a write_filter defense and a seed() bypass for pre-injection attacks. ToolRegistry gates five tools by permission level (READ→ADMIN): parse_pdf, get_webpage, write_memory, read_memory, and send_report. MultiAgentOrchestrator coordinates the two-agent relay in memory mode (Agent A writes; Agent B reads). PropagationLogger matches the canary regex (SECRET-[A-F0-9]{8}) against every tool call, emitting stage-labeled JSONL events. :

- MemoryStore, a key-value store with a write_filter defense and a seed() bypass for pre-injection attacks;
- ToolRegistry, which gates five tools by permission level (from READ to ADMIN): parse_pdf, get_webpage, write_memory, read_memory and send_report;
- MultiAgentOrchestrator, which coordinates the two-agent relay in memory mode (Agent A writes; Agent B reads);
- PropagationLogger, which matches the canary regex SECRET-[A-F0-9]{8} against every tool call and emits stage-labeled JSONL events.

Figure 2 and the tool list above are simplified; scenario-specific tools include query_db (tool_poison), escalate_privilege (permission_esc), and transcribe_audio (audio injection pilot) this tool list are simplified. Several scenarios add their own tools: the tool_poison scenario can also query a database with query_db; in permission_esc, Agent B holds ADMIN tools such as delete_record and escalate_privilege; and the audio pilot transcribes clips with transcribe_audio.

3.2 Kill-Chain Stages and Attack Scenarios

Threat model. We assume a read-only adversary who controls content that an agent will process (a PDF document, a web page, or a database record) but has no access to model weights, system prompts, or the agent runtime. The adversary’s goal is to cause the agent to issue a specific unauthorized outbound action (send_report send_report to an attacker-controlled address) while the legitimate task still succeeds (the dual-completion attack pattern [2]). Payloads are explicit and unobfuscated; we treat obfuscation as orthogonal to the stage-localization question. In the relay scenarios Agent A and Agent B communicate only through the shared MemoryStore




Figure 2: Benchmark setup. Top: the user task and an attacker payload carrying its canary token. Middle: the attack surfaces (named as in Section 3.2), which reach Agent A as tool results, except pre-seeded memory, which enters shared memory directly; the two-agent pipeline (Agent A reads and writes shared memory; Agent B reads memory and calls `send_report`); the four checkpoints (1 Exposed, 2 Persisted, 3 Relayed, 4 Executed); and where each defense looks. Dashed: in the single-agent `tool_poison` scenario the agent that reads the tool result sends the report itself. `memory_poison` is also single-agent (its agent is drawn as Agent B), and `permission_esc` passes the injection from Agent A to Agent B in a delegation message instead of shared memory. Bottom: models, attack surfaces, defense conditions and run counts: 8–36 runs per cell on text surfaces (Table A1), 3 per cell in the PDF relay (Table 4 and Figure 5), and 4 per model for the audio pilot (Table 5).

(in `permission_esc`, through a delegation message); neither agent has direct access to the other’s internal state.

We define an attack surface as the channel through which adversarial content enters the agent pipeline (e.g., pre-seeded memory_{state}, tool response stream, or web document, a tool result, or a web page). Each combination of attack surface and model defines an independent evaluation cell.

Each attacked run is annotated with the furthest kill-chain stage the canary reaches [1]: **EXPOSED** **Exposed** (canary in any tool result), **PERSISTED** **Persisted** (canary in a `write_memory` `write_memory` call), **RELAYED** **Relayed** (canary read by Agent B), **EXECUTED** **Executed** (canary in any outbound tool argument). The gap between **EXPOSED** and **PERSISTED** **Exposed** and **Persisted** identifies summarization-stage filtering; between **PERSISTED** and **EXECUTED** **Persisted** and **Executed**, it identifies execution-stage refusal.

Four text-surface scenarios and one multimodal scenario

cover distinct injection channels: . Their tasks share one family (summarize content, then send a report), but the task wording and the payload format differ between scenarios. The six attack surfaces map to scenarios as follows: web page (propagation, `permission_esc`), pre-seeded memory (`memory_poison`), tool result (`tool_poison`), visible PDF text (`pdf_append`) and white PDF text (`pdf_whitefont`) in the PDF relay, and audio transcript (`audio_inject`, a pilot); multi-surface combines web page and pre-seeded memory, and `pdf_metadata` is a PDF pilot.

`memory_poison` `memory_poison` (pre-seeded memory; single agent) — `MemoryStore.seed()` the harness pre-loads the injection into memory before the run, using `MemoryStore.seed()`. Task: summarize `research_notes` and email to `supervisor@lab.org`. The agent calls `read_memory()` `read_memory`, receives the poisoned record, and

may call `send_report(to="exfil@attacker.com", ...)` `send_report(to="exfil@attacker.com", ...)`.

`tool_poison` `tool_poison` (tool result; single agent) — injection arrives in the return value of `get_webpage()` or `query_db()` `get_webpage` or `query_db`. The legitimate document and the adversarial instruction appear in the same tool result, and the agent that reads it also sends the report.

`propagation` `propagation` (web page; text relay) — two-agent memory relay. Agent A calls `get_webpage()` `get_webpage` on a poisoned page and `write_memory("summary", ...)` writes a summary to memory with `write_memory`. Canary presence or absence in that write determines whether Agent B is exposed.

`permission_esc` `permission_esc` (web page; privilege escalation) — injection calls `escalate_privilege()` before issuing a high-privilege action. Agent A reads a poisoned web page and writes a delegation message for Agent B, which holds ADMIN tools; the payload asks Agent A to have Agent B delete a record named by the canary token. The attack succeeds if Agent B calls an ADMIN tool (`delete_record`, `exfiltrate_data` or `escalate_privilege`). Tests two-step adversarial instruction following.

`cross_modal_relay` (Phase 3 `cross_modal_relay` (visible and white PDF text; PDF relay) — a three-boundary scenario. Agent A calls `parse_pdf()` on parses an injected document and `write_memory("doc_summary", ...)` with `parse_pdf` and writes a summary to memory with `write_memory`. Agent B calls `read_memory()` and `send_report()` `read_memory` and `send_report`. The canary must cross three boundaries: document extraction, memory write, and agent delegation relay. Attack variants: `pdf_append` `pdf_append` (visible appended text), `pdf_whitefont` `pdf_whitefont` (1 pt white-text overlay, human-invisible), `pdf_metadata` `pdf_metadata` (payload in /Keywords and /Subject fields). We also test heterogeneous model cross-model pairs (Agent A $\neq$ Agent B) to measure whether the relay node writer’s identity determines downstream exposure.

3.3 Models, Defenses, and Metrics

Five frontier models We test five production LLMs: GPT-4o-mini (gpt-4o-mini-2024-07-18), GPT-5-mini², DeepSeek Chat v3 (deepseek-chat (v3)), Claude Haiku 4.5, and Claude Sonnet 4.5.² “Claude” below refers to both Claude models. Temperature = 0.0; max tokens = 2048. 950 total runs: 764 text-surface runs across 9 batches (March 10–27, 2026) and 186 Phase 3 multimodal PDF relay runs (April 1, 2026).

²Accessed via the OpenAI API during the experimental period (March 2026) under the identifier gpt-5-mini-2025-02-15; the publicly released model name may differ at time of publication.

³API identifiers: gpt-4o-mini-2024-07-18 for GPT-4o-mini, gpt-5-mini-2025-02-15 for GPT-5-mini, and deepseek-chat for DeepSeek Chat.

Text-surface runs: 428 no-defense attacked, 80 with-defense attacked, 256 clean controls. Phase 3 PDF relay runs: 138 homogeneous pairs (Block A) + in same-model pairs and 48 in cross-model pairs (Block B), $n = 3$ per cell. While text-surface sample sizes are modest (Text-surface cells have $n = 8\text{--}36$), the key findings replicate across independent batches spaced; several of them were collected in independent batches up to 17 days apart. Phase 3 cells at and agree across batches. PDF relay cells have $n = 3$ are pilot-scale; confidence intervals are reported throughout.

Five conditions: none (control baseline); `write_filter` `write_filter` (MemoryStore keyword scan before commit) [9]; `pi_detector` `pi_detector` (secondary LLM classifies each outgoing query); `spotlighting` (tool results wrapped in `<document>` XML delimiters `spotlighting` (untrusted content wrapped in delimiters, with an instruction to treat it as data) [9]; all (all three combined).

ASR = fraction of attacked runs where canary reaches EXECUTED Executed. Task success = fraction of all runs where the legitimate task goal is completed. All fractions are reported with Wilson 95% confidence intervals (CI) as point estimate (95% CI: lo–hi) [10].

4 Results

4.1 Exposure: 100% Among Runs That Called the Tool

Table 1 reports overall ASR attack and task success per model. In our evaluation, every model receives every injection: exposure rate is 100% across all 764 attacked runs. The safety gap is entirely downstream of context exposure. Figure ?? illustrates three qualitatively distinct regimes among runs that called the tool: the payload sits in the result of a tool the agent must call for its task, and all 372 no-defense attacked text-surface runs that made a tool call were exposed. The other 56 of the 428 made no tool call and never read the payload: 48 GPT-5-mini runs (16 each in `memory_poison`, `tool_poison` and `permission_esc`) and 4 runs each of DeepSeek Chat and GPT-4o-mini (`permission_esc`), all from one collection batch (`high_n_v1` in the `batch` column of the released logs). They count as unsuccessful attacks in Table 1 and in Section 4.3; none is in the text relay. GPT-5-mini’s 94% task success and its 48 of 136 attacked runs without a tool call come from different batches: `high_n_v1` holds all 48 (out of its 100 GPT-5-mini attacked runs) and no clean controls, while the 36 GPT-5-mini clean-control runs come from three batches (`scenario_compare_v1`,

²Table 1 includes 16 pilot runs (4 GPT-4o-mini, 8 GPT-5-mini, 4 Claude Sonnet) from audio and `pdf_metadata` injection pilots not broken out in Table A1; see Section 4.9 for details.

³Besides the four scenarios broken out in Section 4.3, Table 1 includes 16 runs of a fifth text scenario, `multi_surface` (web page and pre-seeded memory in one two-agent relay): 4 GPT-4o-mini runs, all executed, and 8 GPT-5-mini and 4 Claude Sonnet 4.5 runs, none executed.

Table 1: Attack and task success per model, text surfaces. n : no-defense attacked runs, including runs that made no tool call. ASR: share of those runs that reached Executed, with Wilson 95% CI. Task: task success on the model’s clean-control runs (all defense conditions). Exposure was 100% among runs that called the tool (Section 4.1).

Model	n	Task	ASR (95% CI)
GPT-4o-mini	60	90%	53% (41–65%)
DeepSeek Chat	68	100%	25% (16–37%)
GPT-5-mini	136	94%	3% (1–7%)
Claude Haiku 4.5	80	100%	0% (0–5%)
Claude Sonnet 4.5	84	100%	0% (0–4%)

Table 2: Kill-chain stage reached, text relay (propagation), no-defense attacked runs. Exposure was 100% among runs that called the tool, and every text relay run called it, so Exposure is 100% in every row.

Model	n	Persisted	Relayed	Executed
GPT-4o-mini	8	100%	100%	100%
DeepSeek Chat	8	100%	100%	100%
GPT-5-mini	20	15%	15%	15%
Claude Haiku 4.5	20	0%	0%	0%
Claude Sonnet 4.5	20	0%	0%	0%

gpt5mini_propagation_v1, multi_surface_v1) in which every GPT-5-mini run called a tool.

Table 1 reports task success next to attack success, the joint view AgentDojo [4] also reports. GPT-4o-mini (90% utility task success, 53% ASR) is the worst case: high task success and high injection compliance. Claude (100% utility, 0% ASR) appears Pareto-efficient within our evaluation—both Claude models achieve completes the task and follows the injection most often. Claude Haiku 4.5 and Claude Sonnet 4.5 reach 100% task success on clean runs while producing and executed none of their 164 text-surface attacks (0% ASR under attack, though we caution this result is specific to /80 and 0/84); this holds for the evaluated scenarios and tasks. DeepSeek only. DeepSeek Chat’s 25% overall ASR masks a combines 0%/100% surface split discussed in 24 and 8/8 on different channels (Section 4.3).

4.2 Where the Canary Token Stops

Table 2 and Figure 3 illustrate per-stage canary survival. Claude eliminates the injection during write_memory: in give the stage reached per run. In the text relay the canary token never appeared in a Claude memory write: 0/40 Claude propagation runs did the canary survive into MemoryStore runs (20 Haiku Claude Haiku 4.5 + 20 Sonnet Claude Sonnet 4.5; 95% CI: 0–8%). The downstream Agent B receives a semantically correct but adversarially clean summary then read a summary that did not contain the token.

The architectural implication is notable: in any multi-agent pipeline, the safety of the relay/summarizer node determines downstream exposure independently of the terminal agent’s safety level. A Claude relay decontaminates for any downstream consumer; A Claude writer kept the canary token out of

shared memory in every relay run we report: the 40 text relay runs here, and the 12 same-model and 3 cross-model PDF relay runs with a Claude writer (Section 4.8). One cross-model pairing (Claude Haiku 4.5 writer, GPT-4o-mini reader, $n = 3$) is consistent with this protecting the reader (Section 4.8); other pairings were not tested. With a GPT-4o-mini relay propagates injections with full fidelity. This maps to or DeepSeek Chat writer, the token reached shared memory in every text relay run (8/8 each). In terms of the Prompt Infection self-replication mechanism [6]: replication is blocked specifically at the write_memory summarization step, the token did not replicate past the Claude memory write. We do not isolate whether this effect arises from training data, system prompts, or tool interaction design; further ablation is required.

For GPT-5-mini partially filters at PERSISTED (the token reached a memory write in 3/20 runs (15%; 95% CI: 5–36%), a non-trivial pass-through rate that rules out both between full compliance and full resistance.

4.3 Attack Success by Surface

Figure 4 reports attack success per model and attack surface; Table A1 and Figure 4 report per-surface ASR. DeepSeek’s result is operationally the most important in Appendix A lists the same text-surface counts. DeepSeek Chat shows the largest swing across scenarios: 0/24 on memory_poison memory_poison (three independent batches spanning 17 days) and 8/8 on tool_poison and propagation. A single-surface evaluation would give DeepSeek tool_poison and propagation. A single-scenario evaluation would have reported either a perfect score (memory-only) or a total failure (tool-stream-only)—either result would be a complete mischaracterization of its actual safety posture.

The most parsimonious or a failing score for the same model. The scenarios differ in task wording and payload format as well as in channel (Section 3.2), so the swing cannot be attributed to the channel alone. One possible explanation is channel-differentiated trust: DeepSeek appears to Chat may treat pre-stored memory records with higher more suspicion than live function-call response streams. This cannot be verified from external behavior, but the pattern replicated cleanly across batches spaced 1–17 days apart. We note this observation is based on black-box behavioral analysis and requires white-box investigation to confirm results.

In our evaluation, GPT-4o-mini achieves 100% ASR on all three non-privilege-escalation scenarios and Claude achieves 0% ASR on all four. The permission_esc permission_esc near-zero results (2/132 total: DeepSeek Chat 1/28 and GPT-5-mini 1/36) most likely reflect payload design—the two-step overtly adversarial payload—rather payload—rather than genuine privilege-escalation resistance. We do not claim this behavior generalizes beyond the evaluated configurations; further validation across additional models and tasks is required.

4.4 Defense Failures

In our evaluation, all four active defense conditions produce 100% ASR on GPT-4o-mini and DeepSeek across propagation and tool_poison ($n = 8$ per cell).




Figure 3: Last stage the canary token reached, as a share of attacked runs per model, no defense. Left: text relay (propagation), $n = 8$ – 20 per model (Table 2). Middle and right: PDF relay, same-model pairs, visible text (pdf_append) and white text (pdf_whitefont), $n = 3$ per model (Table 4). Green: Exposed only (the token was absent from the memory write); amber:elayed to Agent B but not Executed; red: Executed. Hatched: GPT-5-mini never called `parse_pdf`, so it never read the payload; every text relay run called its tool.

Spotlighting reduces On the text relay and on `tool_poison`, no defense condition lowered the attack success of GPT-4o-mini task success from 63% to 50% without reducing ASR—a strictly worse utility-security outcome.

The primary failure mechanism appears to be a mismatch between each defense’s threat model or DeepSeek Chat: all 64 attacked runs under `write_filter`, `pi_detector`, `spotlighting` or all three combined reached Executed (4 runs per model, scenario and defense, all from the defense ablation batch in the released logs, `defense_ablation_v1`). `pi_detector` and our injection surfaces. Spotlighting [9] wraps document content in XML delimiters, but our injection enters via the function-call response stream. `pi_detector` `write_filter` failed on channels they do not inspect. `pi_detector` scans outgoing tool queries for adversarial intent, but our injection arrives payloads arrive in incoming tool results. `write_filter` intercepts `write_filter` scans agent-initiated memory writes, but `memory_poison` pre-seeds the payload via `MemoryStore.seed()` before the agent session begins, bypassing the interceptor entirely.

Each defense is correctly implemented for its stated threat model. The problem is that none of those threat models cover the injection surface under test. This extends the finding `tool_poison` does not route the payload through a memory write, and `memory_poison` pre-seeds it with `MemoryStore.seed()` before the session begins: with `write_filter` active, GPT-4o-mini still executed 8/8 `memory_poison` attacks, as it did without a defense in the same batch (DeepSeek Chat: 0/8 in both conditions). `spotlighting` failed on content it wraps: in the text relay and `tool_poison` the harness places the fetched page or search result that carries the payload inside its delimiters. `write_filter` also failed on the text relay, whose payload passes through the memory write

Table 3: Per-step objective drift (TF-IDF cosine distance from the task description) for one clean and one attacked GPT-4o-mini `memory_poison` run.

Step	Tool	Clean	Attacked	Δ
1	read_memory	0.563	0.571	+0.008
2	send_report (legit)	0.714	0.747	+0.033
3	send_report (harm)	—	0.868	+0.154

it scans, yet blocked the PDF relay (Section 4.8); we cannot explain this difference from our runs. Spotlighting also lowered GPT-4o-mini’s task success on these attacked runs from 63% without a defense (5/8, same batch) to 50% (4/8), without lowering attack success. None of these failures required adapting the attack to the defense, unlike the adaptive attacks of Nasr et al. [?]—who showed adaptive attacks break 12/12 defenses—by demonstrating that non-adaptive attacks can achieve the same result via surface mismatch alone. Any deployment that evaluates defenses only against the surfaces they were designed for will miss every injection that enters from a different channel[8].

4.5 Objective Drift as a Post-Hoc Forensic Signal

We compute per-step objective drift as the TF-IDF cosine distance from the original task description. Table 3 reports step-level drift for GPT-4o-mini on `memory_poison` `memory_poison`: steps 1–2 are statistically indistinguishable nearly identical ($\Delta \leq 0.033$); the signal diverges only at the harmful step itself (Figure A1a in Appendix A).

Last-step drift Δ by scenario: `memory_poison` +0.348, `permission_esc` +At the last step, drift rises by $\Delta = 0.348$ for `memory_poison`, 0.312, `tool_poison` +for `permission_esc`, 0.089, `propagation` +for `tool_poison` and 0.052 for

	Text surfaces (n = 8-36 per cell)				PDF relay (n = 3 per cell)	
GPT-4o-mini	12/12 [76-100%]	8/8 [68-100%]	8/8 [68-100%]	0/28 [0-12%]	0/3 [0-56%]	1/3 [6-79%]
DeepSeek Chat	0/24 [0-14%]	8/8 [68-100%]	8/8 [68-100%]	1/28 [1-18%]	3/3 [44-100%]	3/3 [44-100%]
GPT-5-mini	0/36 [0-10%]	0/36 [0-10%]	3/20 [5-36%]	1/36 [0-14%]	no parse_pdf call (n=3)	no parse_pdf call (n=3)
Claude Haiku 4.5	0/20 [0-16%]	0/20 [0-16%]	0/20 [0-16%]	0/20 [0-16%]	0/3 [0-56%]	0/3 [0-56%]
Claude Sonnet 4.5	0/20 [0-16%]	0/20 [0-16%]	0/20 [0-16%]	0/20 [0-16%]	0/3 [0-56%]	0/3 [0-56%]
	pre-seeded memory memory_poison	tool result tool_poison	web page propagation	web page permission_esc	visible PDF text pdf_append	white PDF text pdf_whitefont

Figure 4: Attack success by model and injection channel, no defense: runs that reached Executed over attacked runs (k/n), with Wilson 95% CIs in brackets. Column headers give the attack surface and the scenario. Text-surface cells ($n = 8-36$, including runs that made no tool call) are also listed in Table A1; PDF relay cells are same-model pairs ($n = 3$), the Executed columns of Table 4. Hatched: GPT-5-mini never called `parse_pdf`. Outlined: DeepSeek Chat’s 0/24 on `memory_poison` and 8/8 on `tool_poison` (Section 4.3).

propagation.

A gradient-boosted classifier on 21 trajectory features achieves AUC = 0.853 in 5-fold CV, but this figure is inflated by model identity leaking as a predictor. The **honest stricter** evaluation—leave-one-scenario-out—collapses to AUC = 0.39–0.57 (chance level) across all held-out scenarios. `objective_drift_mean_after_exposure` accounts for 32.3% of feature importance (3× the next feature; Figure A1b), because it averages drift only across steps after the canary enters context, suppressing the stable pre-injection baseline and amplifying the post-execution spike. That spike arrives concurrent with harm, not before. Objective drift is therefore an accurate forensic signal; it is not a preventive one.

4.6 Provenance Attribution and Dual-Completion

We implement a token-overlap provenance heuristic: for each outgoing tool argument, alphanumeric tokens (≥ 4 chars) are matched against all prior tool results and memory state (threshold: ≥ 3 shared tokens or $\geq 20\%$ argument coverage). **On We applied it to the 22 compromised runs of the first collection batch (scenario_compare_v1: 12 GPT-4o-mini, 9 DeepSeek Chat and 1 GPT-5-mini run, all without a defense); it reconstructed 22/22 (100%) injection paths were correctly reconstructed injection paths with 0 false attributions. Compromised runs from later batches, which make up the rest of the compromised runs in Table 1, were not traced.**

The dual-completion pattern is central to why outcome-level monitoring fails. **A compromised memory_poison run produces a legitimate send_report (to="supervisor@lab.org") In a compromised memory_poison run, the agent first sends the legitimate report to supervisor@lab.org at step 2 (the task succeeds) and a malicious send_report, then a second send_report carrying the canary token at step 3 (the injection executes)—provenance traces both arguments to tool_result:read_memory @step1.; provenance traces the arguments of both calls to the read_memory result at step 1.** Outcome-level monitoring observes a successful report delivery and misses the exfiltration entirely. Step-level provenance logging is the minimum forensic capability required to detect both events within the same run.

4.7 Execution Lag and Detection Feasibility

GPT-4o-mini’s median exposure-to-execution lag is 1 step (mean 1.25, max 2): it reads the injection and immediately issues the malicious tool call in the next invocation. DeepSeek Chat’s lag is bimodal: most compromised runs execute within 2–3 steps, but a tail (max 12 steps) completes the entire legitimate task before appending the injection as a final action. GPT-5-mini when compromised exhibits a mean lag of 3.0 steps.

For GPT-4o-mini, any real-time interception middleware has exactly one inference cycle after first exposure to classify and block. For DeepSeek Chat’s late-executing runs, a mid-run detector calibrated to GPT-4o-mini’s 1-step pattern

will **would** miss the attack entirely. Compromised GPT-4o-mini runs also show excess `send_report` `send_report` calls (mean 1.26 vs. 0.57 clean, $\Delta+0.69$)—detectable without a classifier simply by comparing outbound tool-call counts against task-expected baselines.

4.8 PDF Relay Experiment

To assess whether the kill-chain decomposition generalises beyond text injection, we extend the benchmark to two new attack surfaces (PDF documents, audio transcripts) and introduce the PDF relay experiment introduces a three-boundary kill-chain kill chain, building on the multimodal prompt injection findings work of Ding et al. [?]: **Exposed** [13]: **Exposed** (document extraction), **Persisted** (memory write), **Relayed** (Agent B reads memory), **Executed** (Agent B issues `send_report` `send_report` carrying the canary). We ran 186 trials across five models using the `cross_modal_relay` `cross_modal_relay` scenario ($n = 3$ per cell; 46 homogeneous cells in Block A, cells of same-model pairs and 16 cross-model cells in Block B) cells of cross-model pairs). All relay runs inject through PDF documents (visible and white PDF text); the audio surface was tested only in the pilot of Section 4.9. At $n = 3$, per cell the confidence intervals are wide; all Phase 3 results should be interpreted as directional pilot findings pending larger replication, and we read these results as directional.

Table ?? and Figure ?? report

Three-boundary propagation (same-model pairs). Table 4 reports per-stage rates for homogeneous model same-model pairs under no defense (attacked runs only).

The pattern from text-surface experiments replicates cleanly across ; Figure 3 (middle and right) shows the same runs by last stage reached. The text relay pattern also holds for the PDF surface. DeepSeek propagates Chat propagated every injection through all three boundaries to execution (100% 3/3 per variant; Wilson 95% CI: 44–100%). Claude eliminates every injection at the **PERSISTED** stage With Claude as both agents, the canary token was absent from every memory write (0/6 Haiku Claude Haiku 4.5, 0/6 Sonnet Claude Sonnet 4.5; CI: 0–46%/0–46%), identical to the text-surface results in the text relay. GPT-4o-mini reaches **RELAYED** at 100% but executes **reached Relayed** in every run but executed inconsistently: 0% for `pdf_append` `pdf_append` and 33% for `pdf_whitefont` `pdf_whitefont`.

The invisible-injection finding is notable: `pdf_whitefont` (White text (`pdf_whitefont`, 1 pt white text, invisible to human readers) achieves higher ASR than `pdf_append` on GPT-4o-mini (succeeded at least as often as visible appended text (`pdf_append`): 33% vs. 0%). The whitefont on GPT-4o-mini (1/3 vs. 0/3) and 100% for both on DeepSeek Chat. The white text is extracted first by `pypdf`, placing the directive before the benign abstract in the extracted string; the position bias amplifies compliance. DeepSeek this position may explain the GPT-4o-mini difference. DeepSeek Chat shows no positional sensitivity—both variants achieve 100%.

GPT-5-mini fails to reach **Exposed** never reached **Exposed** (0/6 EXP = 0%). The model did not invoke the `parse_pdf` `parse_pdf` tool in any of the 18 GPT-5-mini runs, producing 0% utility. This is a tool-calling regression: GPT-5-mini-2025-02-15 either does not support our tool schema format or requires a different prompt structure; we report the failure transparently; from black-box runs we cannot tell whether this reflects the model or our tool schema. All GPT-5-mini cells are excluded from further multimodal PDF relay analysis.

`write_filter` achieves 0% EXE universally

Defense effectiveness (same-model pairs). `write_filter` blocked execution in every PDF relay cell tested: 0/3 DeepSeek `pdf_append` Chat `pdf_append`, 0/3 DeepSeek `pdf_whitefont` Chat `pdf_whitefont`, 0/6 GPT-4o-mini `pdf_whitefont`. This extends the text-surface finding: `write_filter` is the only mechanism that blocks at the correct stage (memory write), and it does so across all evaluated injection surfaces `pdf_whitefont`. In the text relay it did not (Section 4.4), although both relays pass the payload through the memory write it scans; we cannot explain the difference from these runs.

`pi_detector` completely fails `pi_detector` did not block on the PDF surface: DeepSeek `pdf_append` Chat `pdf_append` → 100% EXE under `pi_detector` Executed under `pi_detector` (identical to no-defense no defense); GPT-4o-mini `pdf_whitefont` `pdf_whitefont` → 33% (identical to no-defense no defense). This confirms the surface-mismatch hypothesis introduced in is consistent with the channel mismatch of Section 4.4: `pi_detector` `pi_detector` scans outgoing tool queries for adversarial intent, but PDF injection, but the PDF payload arrives in the incoming parsed document content. The defense is correctly implemented for its stated threat model; that threat model simply does not include this surface.

Cross-model pairs. Block B pairs non-identical agents: In cross-model pairs Agent A ($\neq$ Agent B) reads the PDF and writes to shared memory; Agent B reads memory and issues `send_report`. Table ?? and Figure ?? report results for `pdf_append` `send_report`. Figure 5 reports results for `pdf_append`, no defense ($n = 3$ per cell, attacked runs only).

Three directional observations emerge from Table ??; given, with Wilson CIs in its caption. Given $n = 3$ per cell (95% CI ranging from as wide as 21–94%), these should be read as hypotheses for larger-scale follow-up rather than settled findings the three observations below are directional.

Observation 1 (write-time vs. read-time defense). When Claude Haiku is Agent A, PER=0% prevents any downstream propagation. When Claude Haiku is Agent B and inherits pre-written malicious memory, EXE=67% reader position). As a reader, Claude Haiku 4.5 executed 2/3 injections relayed by a GPT-4o-mini writer (95% CI: 21–94%). The directional gap is consistent with a write-stage defense mechanism: the model appears to filter adversarial content when writing to shared memory but not when reading inherited content. Mechanistic validation (e.g., intervention on system prompt vs. RLHF) is required to confirm this interpretation, while the Claude models executed 0/40 raw pre-seeded injections in `memory_poison` (Fig-

Table 4: PDF relay, same-model pairs (Agent A = Agent B), no defense, attacked runs, $n = 3$ per cell. Exposed: canary in the extracted text; Persisted: in the memory write; Relayed: in Agent B’s `read_memory` result; Executed: in Agent B’s `send_report` arguments (= ASR).

Model	pdf_append (visible text)				pdf_whitefont (white text)			
	Exposed	Persisted	Relayed	Executed	Exposed	Persisted	Relayed	Executed
GPT-4o-mini	100%	100%	100%	0%	100%	100%	100%	33%
DeepSeek Chat	100%	100%	100%	100%	100%	100%	100%	100%
Claude Haiku 4.5	100%	0%	0%	0%	100%	0%	0%	0%
Claude Sonnet 4.5	100%	0%	0%	0%	100%	0%	0%	0%
GPT-5-mini	0%	0%	0%	0%	0%	0%	0%	0%

Agent A (reads PDF, writes memory)	GPT-4o-mini	Executed 0/3	Executed 3/3	Executed 2/3
		Persisted 3/3	Persisted 3/3	Persisted 3/3
	DeepSeek Chat	Executed 3/3	Executed 3/3	not tested
		Persisted 3/3	Persisted 3/3	
	Claude Haiku 4.5	Executed 0/3	not tested	Executed 0/3
		Persisted 0/3		Persisted 0/3
	Agent B (reads memory, acts)			
	GPT-4o-mini	DeepSeek Chat	Claude Haiku 4.5	

Figure 5: PDF relay by writer (Agent A, rows) and reader (Agent B, columns): runs that reached Executed and runs whose memory write contained the token (Persisted), out of $n = 3$ per cell; pdf_append, no defense. Diagonal cells are same-model pairs (Table 4). Off-diagonal cells are cross-model pairs, in which Relayed equals Persisted in every pair; Executed with Wilson 95% CIs: GPT-4o-mini writer and Claude Haiku 4.5 reader 2/3 (21–94%); GPT-4o-mini writer and DeepSeek Chat reader 3/3 (44–100%); Claude Haiku 4.5 writer and GPT-4o-mini reader 0/3 (0–56%); DeepSeek Chat writer and GPT-4o-mini reader 3/3 (44–100%). Hatched: pairing not tested.

ure 4). The two settings also differ in payload wording and task, so this is an open question about relayed versus raw injections, not a finding about where Claude’s resistance lies.

Observation 2 (relay node as primary chokepoint). Across both models tested as Agent A, the relay node identity determines whether any canary enters memory at all. A Claude relay (Agent A) produces PER=writer position). With Claude Haiku 4.5 as the writer and GPT-4o-mini as the reader, the token was absent from shared memory and the reader executed 0% regardless of Agent B; a /3. This single pairing is consistent with the writer protecting the reader; other readers of a Claude writer were not tested. A GPT-4o-mini or DeepSeek relay produces PER=Chat writer put the

Table 5: Pilot surface results (no-defense attacked runs, $n = 4$ per cell). These runs are not part of Table 1 or Figure 4.

Surface	Model	n	ASR (95% CI)
audio_inject	GPT-4o-mini	4	0% (0–49%)
audio_inject	Claude Sonnet 4.5	4	0% (0–49%)
pdf_metadata	GPT-4o-mini	4	0% (0–49%)
pdf_metadata	GPT-5-mini	4	0% (0–49%)
pdf_metadata	Claude Sonnet 4.5	4	0% (0–49%)

token into memory in every run (Persisted 100%, shifting all safety burden), leaving the outcome to Agent B. This extends the text-surface result from Section 4.2 to the PDF surface.

Observation 3 (unexplained REL–EXE gap for homogeneous Relayed–Executed gap for same-model GPT-4o-mini). In Block A, With GPT-4o-mini reaches RELAYED=100% for pdf_append but EXE=0% as both agents, pdf_append reaches Relayed in every run but Executed in none: Agent B reads the canary but does not include it in send_report. When DeepSeek is Agent A (Block B) send_report. With DeepSeek Chat as the writer, GPT-4o-mini as Agent B achieves EXE=100%, the reader executed 3/3. A plausible explanation is that DeepSeek Chat writes the directive verbatim into memory while GPT-4o-mini paraphrases it, reducing instruction explicitness below the compliance threshold. This REL–EXE gap is an open empirical question: instruction explicitness in shared memory may be a vulnerability axis independent of canary token presence, but ; the mechanism is unconfirmed at this sample size.

4.9 Audio Injection and PDF Metadata Pilots

In addition to the text and PDF surfaces reported above, we conducted small-scale pilots on two additional injection channels: audio_inject audio_inject (adversarial instructions embedded in a synthesized speech clip processed by transcribe_audio) and pdf_metadata transcribe_audio) and pdf_metadata (payload placed in the PDF /Keywords and /Subject fields, extracted by pypdf).

Table 5 summarizes these results. Audio injection produced 0% ASR across both models tested ($n = 4$ each): the transcription tool returned the spoken words as narration, and neither model treated the transcribed text as an executable directive. This null result is consistent with prior multimodal

injection findings [2] [13] and suggests that audio-to-text transcription may implicitly contextualize content as data rather than instruction. PDF metadata injection also produced 0% ASR ($n = 4$ per model, three models): pypdf extracts /Keywords and /Subject fields, but models did not parse these metadata strings as actionable directives. Both pilot surfaces yielded null results at this sample size; we include them for completeness and to account for the 16-run discrepancy between Table 1 ($n = 428$) and Table A1 ($n = 412$). Larger-scale evaluation with stronger payload designs (e.g., multi-turn audio prompts or metadata fields mimicking system instructions) is needed before drawing conclusions about the safety of these channels.

5 Discussion

Claude’s 0/

Where the canary token stopped. Claude executed none of its 164 ASR (0 successes across all four text-surface scenarios, attacks (no defenses applied) reflects safety behavior that activates at the summarization stage rather than the execution stage. Phase 3 confirms this pattern on the PDF surface, and in both relays the canary token was absent from Claude’s memory writes: 0/6 Claude Haiku 40 text relay runs and 0/6 Claude Sonnet pdf_append 12 same-model PDF relay runs (Table 4). As a reader, however, Claude Haiku 4.5 executed 2/pdf_whitefont runs produced canary persistence (EXP=100%, PER=3 injections relayed by GPT-4o-mini, against 0%). Critically, Block B reveals that this defense is write-time only: when Claude Haiku inherits pre-written malicious memory as Agent B, it propagates at 67%. This is architecturally different from execution-stage refusal: the model distinguishes between processing adversarial content and writing it versus receiving it from a trusted delegation channel. None of the evaluated defense mechanisms reproduced this behavior on any other model. We caution that this finding is specific /40 raw pre-seeded injections for the Claude models. Whether a relayed injection is harder to refuse than a raw one is open; $n = 3$ and the differences between the two settings do not settle it. These results apply to our benchmark and cannot be generalized without further validation only.

Placing a Claude model at the summarization/relay position (Agent A) produces 0/46 canaries entering shared memory across text and PDF surfaces in our evaluation— independent of which model serves as Agent B. Block B provides directional evidence for an important asymmetry: placing Claude as Agent B does not provide the same protection (67% propagation when

Writer position as a candidate design lever. A Claude writer kept the token out of shared memory in every relay run we report, and in the one cross-model pairing tested (GPT-4o-mini writes the malicious memory, though reader, $n = 3$ makes this a hypothesis rather than a settled result). The implication for system design is direct: multi-agent frameworks should expose write-node identity as a first-class architectural choice. Operators can harden pipelines today by requiring safety-certified models at every) the reader then executed nothing. If this holds beyond our settings, the model placed at each inter-agent write boundary, independent of which models are used matters independently of the models elsewhere in the chain. We propose relay decontamination rate at the write stage as a mandatory metric for multi-agent security benchmarks suggest

reporting, as a separate metric, the rate at which the canary token is absent from memory writes.

DeepSeek

Apparent safety depends on the channels evaluated. DeepSeek Chat’s 0%/100% ASR split, pi_detector’s complete 24 vs. 8/8 swing, pi_detector’s failure on PDF injection, and the audio injection null result collectively null result show that a model’s or defense’s apparent safety posture is almost entirely determined by which surfaces depends on which channels an evaluation covers. This is not a fixable implementation detail—it is a consequence of defenses being designed for specific threat models. The correct response is not to patch individual defenses but to change evaluation requirements: any claimed safety result must enumerate every injection surface in the target deployment and demonstrate coverage. Defenses that pass surface-specific evaluations but fail Because the DeepSeek Chat scenarios also differ in task and payload format, that swing is an observation rather than an isolated channel effect. Claimed safety results should enumerate the injection channels they cover; detectors that work within one scenario but not on held-out surfaces ones (AUC 0.39–0.57 for our drift classifiers) should be reported as surface-local, not general scenario-local.

pdf_whitefont equaling or exceeding pdf_append in ASR (100% on DeepSeek, 33%

Parser-level trust for document-handling agents. pdf_whitefont succeeded at least as often as pdf_append (3/3 each on DeepSeek Chat; 1/3 vs. 0%/3 on GPT-4o-mini) shows that the vulnerability is in the extraction pipeline, not, so the payload need not be visible in the rendered document. Any defense that operates on Screening the rendered or displayed PDF—visual PDF (visual screening, human review, content moderation of images—will miss injections that exist only in images) would not see these payloads; a check has to operate on the parser’s extracted text. Future document-handling agents need extraction-level content integrity: cryptographic attestation of source, parse-time payload detection, or sandboxed extraction environments that can flag content-display divergence before it enters agent context.

6 Broader Impact

This work studies adversarial vulnerabilities in LLM agents. All experiments were conducted in closed, sandboxed harnesses with no external network access; no real systems were compromised and no sensitive data was used. The injection payloads and attack surfaces described are representative of techniques already documented in the public literature [2, 3]. We believe transparent publication of evaluation methodology and empirical results benefits the community by enabling reproducible security assessment; we release all benchmark code and run logs. The primary dual-use risk is that surface-mismatch channel-mismatch analysis could inform more targeted attack design. However, we assess that the defensive contribution—establishing that evaluation coverage determines apparent security posture—outweighs this risk,

and that the attack techniques described are already publicly known.

7 Limitations

Sample size. Per-cell n ranges from 3 (Phase 3 PDF relay) to 36 (text-surface scenarios). The strongest text-surface findings results (Claude 0/20 per scenario, DeepSeek memory Chat memory_poison 0/24, GPT-4o-mini memory memory_poison 12/12) are consistent across independent batches; the Phase 3 PDF relay cells at $n = 3$ have wide Wilson CIs (21–94% for the 67% relay result) and must be treated as pilot-scale hypotheses (Claude Haiku 4.5’s 2/3 as a reader).

Single task family. All scenarios use one task type (document summarization with an email delivery action) family (summarize content, then send a report), with task wording and payload format that differ between scenarios. Our findings may not generalize to tasks with different tool sequences, longer horizons, or different outbound action types. This is the primary threat to external validity; replication across diverse task distributions is required before deployment recommendations can be made with confidence.

What the canary measures. The canary token tracks copying of the payload marker. A paraphrased instruction could survive a memory write without it, so a token absent from memory does not show that the instruction is absent; Observation 3 in Section 4.8 is one case where this may matter.

Canary detection confound. A model that refuses an injection but quotes the canary token in its refusal explanation would be incorrectly labeled EXE=True. We inspected all Executed. We manually inspected the 22 compromised text-surface runs manually runs of the first batch (the runs traced in Section 4.6) and found no such false positives, but; compromised runs from later batches were not inspected, and this confound is possible in principle and is not systematically controlled.

Run provenance. The columns of Table 1 are computed over different run sets (task success over clean-control runs, ASR over no-defense attacked runs), and each model’s runs come from a different mix of the collection batches (Section 3).

Explicit payloads. Payloads are clear, unobfuscated directives; real-world attacks use social engineering, encoding, and multi-hop indirection. Surface-mismatch Channel-mismatch results are likely conservative: defenses designed for the wrong surface will channel would also fail against obfuscated off-surface off-channel attacks.

Other limitations. Five frontier models production LLMs from three providers; reasoning models (o3, o4-mini) and open-weight models (Llama 3.3, Mistral) are not evaluated. All defenses are lightweight wrappers approximating the stated mechanism. We do not isolate whether the absence of the token from Claude’s summarization-stage filtering memory writes arises from training, system prompt, or tool API design. permission_esc permission_esc near-zero ASR

(2/132) most likely reflects payload design rather than genuine privilege-escalation resistance. GPT-5-mini-2025-02-15 produced GPT-5-mini produced a 0% tool-call rate in Phase 3 the PDF relay; we cannot determine whether this is a capability regression or API schema incompatibility from black-box evaluation.

Open questions. Do relayed injections get past readers that refuse raw ones? Does a Claude writer protect readers other than GPT-4o-mini? Why did write_filter block the PDF relay but not the text relay? Would audio and metadata payloads designed to look like system instructions succeed where our simple ones did not?

8 Conclusion

Stage-level canary tracking across 950 runs reveals a simple but consequential structure: prompt injection is not a model-capability problem, it is a pipeline-architecture problem. Every model is fully exposed; the safety gap lives entirely in what happens to an injection after it enters context separates where an injection stops from whether it succeeds. Exposure was 100% among runs that called the tool; the outcomes differ downstream. Claude Haiku 4.5 and Claude Sonnet 4.5 executed none of their 164 text-surface attacks, and in the text relay the canary token never appeared in a memory write (0/40); GPT-4o-mini executed 53% of its attacks. We draw four findings and, more importantly, four prescriptive claims about how agentic systems should be built and evaluated.

Claude eliminates injections at write_memory summarization on every surface we tested (0/164 text-surface, 0/

Finding 1: writer position. A Claude writer kept the canary token out of shared memory in every relay run we report (40 text relay, 12 Phase same-model and 3 PDF runs as Agent A; 95% CI 0–2%). This is not a property of Claude’s overall capability level—it is a property of its write-stage behavior. The practical implication is immediate: in any multi-agent pipeline where components can be chosen, routing all inter-agent memory writes through a safety-verified write node provides pipeline-wide decontamination independent of downstream agent choices. Conversely, Block B results show Claude as Agent B propagates at 67%—so write-node placement is the only position where this protection materializes. Future agent frameworks should expose write-node identity as a first-class architectural choice, not a deployment afterthought (cross-model PDF relay runs). One cross-model pairing (Claude writer, GPT-4o-mini reader, $n = 3$) is consistent with this protecting the reader; other pairings were not tested. Whether writer placement protects readers in general is open.

All four evaluated defenses—write_filter, pi_detector, spotlighting, and their combination—fail on at least one injection surface, not because they are incorrectly implemented but because they were designed for a different surface than the one under attack. This failure mode requires no adversarial adaptation. The correct design principle is surface decomposition before defense selection: enumerate every channel through which untrusted content can enter the pipeline, then select or compose defenses that cover each channel’s specific injection point. A monolithic defense evaluated only on the surface it was designed for is not a security evaluation.

The write-vs-read safety asymmetry (Claude blocks at write, propagates at read) suggests that agent memory currently lacks the infrastructure to carry trust provenance: information about

Finding 2: reader position. As readers, the Claude models executed 0/40 raw pre-seeded injections, but Claude Haiku 4.5 executed 2/3 injections relayed by GPT-4o-mini. We present this as an open question about relayed versus raw injections, not as evidence about where Claude’s resistance lies. If relayed content is trusted differently, agent memory would benefit from provenance: which agent wrote a memory record, under what safety context, and from what sourcesurface. If Agent B could interrogate “was this memory written by a safety-certified node from a document I trust?” it could apply calibrated skepticism to inherited content rather than treating all memory reads as equivalently trusted delegation. Implementing content-addressedrecord, provenance-tagged memory stores is a concrete near-term architectural direction for multi-agent frameworksfrom which source.

DeepSeek’s

Finding 3: channel dependence. DeepSeek Chat went from 0%/100% ASR split across memory 24 on pre-seeded memory to 8/8 on tool results; the scenarios also differ in task and payload format, so this is an observed swing, not an isolated channel effect. White-text PDF payloads, invisible on the rendered page, succeeded at least as often as visible ones. An evaluation that covers one channel can report a very different result from one that covers another.

Finding 4: defense coverage. pi_detector and tool-stream surfaces, and the audio injection null result (spoken instructions treated as narration, not directives), show that apparent security posture is determined almost entirely by which surfaces an evaluation covers. A model that achieves 0% ASR on a single surface can achieve 100% on the next. The community standard for agentic security evaluation must require: write_filter failed on channels they do not inspect, and spotlighting failed on content it wraps (the text relay and tool_poison), all without adaptive attacks. write_filter blocked the PDF relay but not the text relay, although both pass through the memory write it scans; this difference is unexplained. We suggest that agent security evaluations report (i) kill-chain stage decomposition stages rather than outcome-only ASR; (ii) multi-surface injection coverage multiple injection channels, including document extraction, tool response streamsresults, audio transcripts, and pre-seeded memory; (iii) cross-model relay testing with heterogeneous agent pairs; and (iv) relay decontamination rate measured separately at results at the write and read positions. Until these become baseline requirements, security evaluations of agentic systems will systematically underestimate real attack surfaces.

Broader practical impact. As institutional NLP pipelines over financial documents—earnings calls, SEC filings, analyst reports—adopt LLM-agent architectures, they inherit the injection surfaces and propagation dynamics characterized above. Write-node placement, surface-aware defense composition, and memory provenance provide actionable architectural levers for securing these deploymentsseparately.

References

- [1] E. M. Hutchins, M. J. Cloppert, and R. M. Amin. Intelligence-driven computer network defense informed by analysis of adversary campaigns and intrusion kill chains. *Proc. 6th Annual International Conference on Information Warfare and Security*, 2011.
- [2] K. Greshake et al. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injections. *AISeC @ CCS*, 2023.
- [3] F. Perez and I. Ribeiro. Ignore previous prompt: Attack techniques for language models. *NeurIPS ML Safety Workshop*, 2022.
- [4] E. Debenedetti et al. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. *NeurIPS Datasets and Benchmarks Track*, 2024.
- [5] Q. Zhan, Z. Liang, Z. Ying, and D. Kang. InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents. *Findings of ACL*, 2024.
- [6] D. Lee and M. Tiwari. Prompt infection: LLM-to-LLM prompt injection within multi-agent systems. *arXiv:2410.07283*, 2024.
- [7] X. Yang, Y. He, S. Ji, B. Hooi, and J. S. Dong. Zombie agents: Persistent control of self-evolving LLM agents via self-reinforcing injections. *Lifelong Agent Workshop @ ICLR*, 2026. *arXiv:2602.15654*.
- [8] M. Nasr, N. Carlini, C. Sitawarin, S. V. Schulhoff, J. Hayes, M. Ilie, et al. The attacker moves second: Stronger adaptive attacks bypass defenses against LLM jailbreaks and prompt injections. *arXiv:2510.09023*, 2025.
- [9] K. Hines et al. Defending against indirect prompt injection attacks with spotlighting. *arXiv:2403.14720*, 2024.
- [10] E. B. Wilson. Probable inference, the law of succession, and statistical inference. *JASA*, 22(158):209–212, 1927.
- [11] Y. Wang, W. Zou, R. Geng, and J. Jia. AgentWatcher: A rule-based prompt injection monitor. *arXiv:2604.01194*, 2026.
- [12] C. Xiang, D. Zagieboylo, S. Ghosh et al. Architecting secure AI agents: Perspectives on system-level defenses against indirect prompt injection attacks. *arXiv:2603.30016*, 2026.
- [13] M. Ding, S. Xia, C. Kong, and X. Jiang. Covert visual prompt injection against commercial multimodal large language models. *arXiv:2603.29418*, 2026.

- [14] Q. Zhang, L. Fu, L. Lian et al. Evaluating privilege usage of agents with real-world tools. *FSE Ideas, Visions and Reflections Track*, 2026. arXiv:2603.28166.
- [15] X. Wang, Y. Zhou, Q. Wang et al. Beyond content safety: Real-time monitoring for reasoning vulnerabilities in large language models. *arXiv:2603.25412*, 2026.
- [16] A. Lynch. The persistent vulnerability of aligned AI systems. PhD thesis, University College London, 2025. arXiv:2604.00324.

A Supplementary Results

Table A1 lists the text-surface counts of Figure 4, and Figure A1 collects the objective-drift results of Section 4.5.

Table A1: Attack success k/n per model and text scenario (no-defense attacked runs, including runs that made no tool call; Figure 4 adds Wilson 95% CIs). The `permission_esc` total of 2/132 comprises DeepSeek Chat 1/28 and GPT-5-mini 1/36.

Model	pre-seeded memory memory_poison	tool result tool_poison	web page propagation	web page permission_esc
GPT-4o-mini	12/12	8/8	8/8	0/28
DeepSeek Chat	0/24	8/8	8/8	1/28
GPT-5-mini	0/36	0/36	3/20	1/36
Claude Haiku 4.5	0/20	0/20	0/20	0/20
Claude Sonnet 4.5	0/20	0/20	0/20	0/20




Figure A1: Objective drift (Section 4.5). (a) Per-step drift for one clean and one attacked GPT-4o-mini `memory_poison` run; values are Table 3; shaded: the harmful `send_report` call. (b) Feature importances of the gradient-boosted classifier (21 features, 5-fold CV, AUC 0.853; the 14 largest shown).